

- [Technical Report: Stateless Proxy Upload Gateway](#)
 - [1. Executive Summary](#)
 - [2. Problem Statement](#)
 - [3. Solution Architecture](#)
 - [3.1 Design Pattern: Time-Partitioned State Machine](#)
 - [3.2 System Components](#)
 - [4. Technical Implementation Details](#)
 - [4.1 Ingestion Flow](#)
 - [4.2 Background Processing](#)
 - [4.3 Support & Recovery](#)
 - [5. Deployment & Configuration](#)
 - [6. Conclusion](#)

Technical Report: Stateless Proxy Upload Gateway

1. Executive Summary

Objective: Enable high-volume file ingestion with zero client-side orchestration, ensuring robustness, scalability, and data consistency.

Solution: This project implements a **Stateless Proxy Upload Gateway** using a **Time-Partitioned State Machine**. The system accepts raw file uploads, automatically generates metadata, and orchestrates downstream processing asynchronously using **MinIO** (Object Storage) and **Kafka** (Messaging).

Key Outcomes:

- **Scalability:** Stateless ingestion tier allows horizontal scaling to handle increasing load.
- **Reliability:** Decoupled ingestion and processing layers prevent backpressure from affecting uploads.
- **Simplicity:** Clients only need to upload files; the server handles batching, metadata generation, and recovery.
- **Maintainability:** No SQL database required for state management; "State" is derived from the file system structure (S3/MinIO) and time.

2. Problem Statement

Traditional file ingestion systems often suffer from:

1. **Client Complexity:** Clients must manage batch IDs, generate metadata, and handle partial failure retries.
2. **State Management Overhead:** maintaining a database to track file status adds latency and operational complexity.
3. **Tight Coupling:** Synchronous processing of uploads blocks clients and creates bottlenecks.

3. Solution Architecture

3.1 Design Pattern: Time-Partitioned State Machine

Instead of a database row tracking every file, we use the storage system itself as the source of truth.

- **State = Directory Location:** A file's location determines its state.
 - **tmp-bucket/data/{batchId}/:** **Ingesting**
 - **tmp-bucket/ready-to-process/{yyyy}/{MM}/{dd}/{HH}/{batchId}:** **Ready for Processing**
 - **prod-bucket/data/{batchId}/:** **Processed & Live**
- **Time Partitioning:** "Ready" markers are organized by hour. This allows the processor to scan small, distinct windows of time rather than the entire bucket, ensuring operations remain $O(1)$ relative to total dataset size.

3.2 System Components

1. Ingestion Service (Spring Boot):

- Exposes **POST /api/batches/upload**.
- Generates a unique **Batch ID (UUID)**.

- Uploads files to MinIO **tmp-bucket**.
- **Auto-Generates Metadata:** Creates companion JSON files for routing.
- Writes an "implicit completion marker" to the **ready-to-process** path.

2. Storage Layer (MinIO):

- **Temporary Bucket:** Holds raw uploads. configured with a **7-Day TTL** Lifecycle Policy to auto-expire old data.
- **Production Bucket:** Long-term storage for processed files.

3. Batch Processor (Background Worker):

- **Schedule:** Runs every 5 minutes (configurable).
- **Logic:** Scans the *current* and *previous* hour's **ready-to-process** prefixes.
- **Action:** Atomically promotes files to **prod-bucket** and publishes metadata events to **Kafka**.

4. Messaging Layer (Kafka):

- **Topic Alpha:** Receives ***-meta1.json** events.
- **Topic Beta:** Receives ***-meta2.json** events.
- Decouples file movement from downstream business logic (e.g., ETL jobs, notifications).

4. Technical Implementation Details

4.1 Ingestion Flow

The client performs a single **multipart/form-data** POST.

```
// IngestionService.java
String batchId = UUID.randomUUID().toString();
// 1. Upload Original File
minioClient.putObject(... objectName ...);
// 2. Generate Metadata
minioClient.putObject(... "meta1.json" ...);
minioClient.putObject(... "meta2.json" ...);
// 3. Mark Ready
minioClient.putObject(... "ready-to-process/2026/01/21/10/" + batchId ...);
```

4.2 Background Processing

The `BatchProcessor` ensures *at-least-once* delivery. It is idempotent:

```
// BatchProcessor.java
public void processTimeWindow(ZonedDateTime time) {
    String prefix = "ready-to-process/" +
time.format(DateTimeFormatter.ofPattern("yyyy/MM/dd/HH")));
    // List all batches in this hour
    for (Result<Item> batch : minioClient.listObjects(prefix)) {
        // Check if already processed (check prod-bucket)
        if (isProcessed(batchId)) continue;

        // Copy to Prod & Send to Kafka
        processBatch(batchId);
    }
}
```

4.3 Support & Recovery

A manual API exists to re-trigger processing for any time range, useful for incident recovery or bug fixes.

- **Endpoint:** `POST /api/batches/reprocess?start=...&end=...`
- **Mechanism:** Re-uses the exact same window scanning logic as the background job.

5. Deployment & Configuration

The solution is fully containerized using Docker Compose.

- **MinIO:** `minio/minio:RELEASE.2023-01-31T02-24-40Z`
- **Kafka:** `confluentinc/cp-kafka:7.3.0`
- **Application:** Java 17 / Spring Boot 3.0.0

Network Resilience: Kafka is configured with `KAFKA_ADVERTISED_LISTENERS` supporting distinct internal (Docker) and external (Host) access, facilitating easy local debugging and monitoring.

6. Conclusion

This architecture delivers a robust, high-throughput ingestion system. By removing state management from the application layer and leveraging the distinct capabilities of Object Storage and Kafka, we achieved a solution that is easy to scale, simple to operate, and resilient to failure.